

Computational Finance and FinTech Numerical and Computational Foundations

Contents

2 Numerical and Computational Foundations	1
2.1 Arrays with Python lists	1
2.2 NumPy arrays	2
2.3 Structured NumPy arrays	9
2.4 Data Analysis with pandas: DataFrame	10

2 Numerical and Computational Foundations

- Further reading: **Py4Fi, Chapters 4 and 5**

2.1 Arrays with Python lists

Introduction to Python arrays

- Before introducing more sophisticated objects for data storage, let's take a look at the built-in Python `list` object.
- A `list` object is a one-dimensional array:

```
[1]: v = [0.5, 0.75, 1.0, 1.5, 2.0]
```

- `list` objects can contain arbitrary objects.
- In particular, a `list` can contain other `list` objects, creating two- or higher-dimensional arrays:

```
[2]: m = [v, v, v]
      m
```

```
[2]: [[0.5, 0.75, 1.0, 1.5, 2.0],
      [0.5, 0.75, 1.0, 1.5, 2.0],
      [0.5, 0.75, 1.0, 1.5, 2.0]]
```

`list` objects

```
[3]: m[1]
```

```
[3]: [0.5, 0.75, 1.0, 1.5, 2.0]
```

```
[4]: m[1][0]
```

```
[4]: 0.5
```

- Feel free to push this to higher dimensions...

```
[5]: v1 = [0.5, 1.5]
      v2 = [1, 2]
```

```
m = [v1, v2]
c = [m, m]
c
```

```
[5]: [[0.5, 1.5], [1, 2]], [[0.5, 1.5], [1, 2]]
```

```
[6]: c[1][1][0]
```

```
[6]: 1
```

Reference pointers

- Important: `list`'s work with **reference pointers**.
- Internally, when creating new objects out of existing objects, only pointers to the objects are copied, not the data!

```
[7]: v = [0.5, 0.75, 1.0, 1.5, 2.0]
m = [v, v, v]
m
```

```
[7]: [[0.5, 0.75, 1.0, 1.5, 2.0],
      [0.5, 0.75, 1.0, 1.5, 2.0],
      [0.5, 0.75, 1.0, 1.5, 2.0]]
```

```
[8]: v[0] = 'Python'
m
```

```
[8]: [['Python', 0.75, 1.0, 1.5, 2.0],
      ['Python', 0.75, 1.0, 1.5, 2.0],
      ['Python', 0.75, 1.0, 1.5, 2.0]]
```

Python array class

- Python also has an `array` module
- See [Documentation](#)

2.2 NumPy arrays

NumPy arrays

- NumPy is a library for richer array data structures.
- The basic object is `ndarray`, which comes in two flavours:

Object type	Meaning	Used for
<code>ndarray</code> (regular)	<i>n</i> -dimensional array object	Large arrays of numerical data
<code>ndarray</code> (record)	2-dimensional array object	Tabular data organized in columns

`ndarray`

Source: Python for Finance, 2nd ed.

- The `ndarray` object is more specialised than the `list` object, but comes with more functionality.
- An array object represents a multidimensional, homogeneous array of fixed-size items.
- Here is a useful [tutorial](#)

Regular NumPy arrays

- Creating an array:

```
[9]: import numpy as np # import numpy  
a = np.array([0, 0.5, 1, 1.5, 2]) # array(...) is the constructor for ndarray's
```

```
[10]: type(a)
```

```
[10]: numpy.ndarray
```

- ndarray assumes objects of the same type and will modify types accordingly:

```
[11]: b = np.array([0, 'test'])  
b
```

```
[11]: array(['0', 'test'], dtype='<U21')
```

```
[12]: type(b[0])
```

```
[12]: numpy.str_
```

Constructing arrays by specifying a range

- `np.arange()` creates an array spanning a range of numbers (= a sequence).
- Basic syntax: `np.arange(start, stop, steps)`
- It is possible to specify the data type (e.g. `float`)
- To invoke an explanation of `np.arange` (or any other object or method), type `np.arange?`

```
[13]: np.arange?
```

```
[14]: np.arange(0, 2.5, 0.5)
```

```
[14]: array([0. , 0.5, 1. , 1.5, 2. ])
```

NOTE: The interval specification refers to a half-open interval: `[start, stop)`.

ndarray methods

- The `ndarray` object has a multitude of useful built-in methods, e.g.
 - `sum()` (the sum),
 - `std()` (the standard deviation),
 - `cumsum()` (the cumulative sum).
- Type `a.` and hit `TAB` to obtain a list of the available functions.
- More documentation is found [here](#).

```
[15]: a.sum()
```

```
[15]: 5.0
```

```
[16]: a.std()
```

```
[16]: 0.7071067811865476
```

```
[17]: a.cumsum()
```

```
[17]: array([0. , 0.5, 1.5, 3. , 5. ])
```

Slicing 1d-Arrays

- With one-dimensional `ndarray` objects, indexing works as usual.

```
[18]: a[1]
```

```
[18]: 0.5
```

```
[19]: a[:2]
```

```
[19]: array([0. , 0.5])
```

```
[20]: a[2:]
```

```
[20]: array([1. , 1.5, 2. ])
```

Mathematical operations

- Mathematical operations are applied in a **vectorised** way on an `ndarray` object.
- Note that these operations work differently on `list` objects.

```
[21]: l = [0, 0.5, 1, 1.5, 2]  
l
```

```
[21]: [0, 0.5, 1, 1.5, 2]
```

```
[22]: 2 * l
```

```
[22]: [0, 0.5, 1, 1.5, 2, 0, 0.5, 1, 1.5, 2]
```

- `ndarray`:

```
[23]: a = np.arange(0, 7, 1)  
a
```

```
[23]: array([0, 1, 2, 3, 4, 5, 6])
```

```
[24]: 2 * a
```

```
[24]: array([ 0,  2,  4,  6,  8, 10, 12])
```

Mathematical operations (cont'd)

```
[25]: a + a
```

```
[25]: array([ 0,  2,  4,  6,  8, 10, 12])
```

```
[26]: a ** 2
```

```
[26]: array([ 0,  1,  4,  9, 16, 25, 36])
```

```
[27]: 2 ** a
```

```
[27]: array([ 1,  2,  4,  8, 16, 32, 64])
```

```
[28]: a ** a
```

```
[28]: array([ 1,  1,  4, 27, 256, 3125, 46656])
```

Universal functions in NumPy

- A number of universal functions in NumPy are applied element-wise to arrays:

```
[29]: np.exp(a)
```

```
[29]: array([ 1.          ,  2.71828183,  7.3890561 , 20.08553692,  
          54.59815003, 148.4131591 , 403.42879349])
```

```
[30]: np.sqrt(a)
```

```
[30]: array([0.          , 1.          , 1.41421356, 1.73205081, 2.          ,  
          2.23606798, 2.44948974])
```

Multiple dimensions

- All features introduced so far carry over to multiple dimensions.
- An array with two rows:

```
[31]: b = np.array([a, 2 * a])  
b
```

```
[31]: array([[ 0,  1,  2,  3,  4,  5,  6],  
          [ 0,  2,  4,  6,  8, 10, 12]])
```

- Selecting the first row, a particular element, a column:

```
[32]: b[0]
```

```
[32]: array([0, 1, 2, 3, 4, 5, 6])
```

```
[33]: b[1,1]
```

```
[33]: 2
```

```
[34]: b[:,1]
```

```
[34]: array([1, 2])
```

Multiple dimensions

- Calculating the sum of all elements, column-wise and row-wise:

```
[35]: b.sum()
```

```
[35]: 63
```

```
[36]: b.sum(axis = 0)
```

```
[36]: array([ 0,  3,  6,  9, 12, 15, 18])
```

```
[37]: b.sum(axis = 1)
```

```
[37]: array([21, 42])
```

Note: `axis = 0` refers to column-wise and `axis = 1` to row-wise.

Further methods for creating arrays

- Often, we want to create an array and populate it later.
- Here are some methods for this:

```
[38]: np.zeros((2,3), dtype = 'i') # array with two rows and three columns
```

```
[38]: array([[0, 0, 0],  
          [0, 0, 0]], dtype=int32)
```

```
[39]: np.ones((2,3,4), dtype = 'i') # array dimensions: 2 x 3 x 4
```

```
[39]: array([[[1, 1, 1, 1],  
            [1, 1, 1, 1],  
            [1, 1, 1, 1]],  
          [[1, 1, 1, 1],  
            [1, 1, 1, 1],  
            [1, 1, 1, 1]]], dtype=int32)
```

```
[40]: np.empty((2,3))
```

```
[40]: array([[1.          , 1.41421356, 1.73205081],  
          [2.          , 2.23606798, 2.44948974]])
```

Further methods for creating arrays

```
[41]: np.eye(3)
```

```
[41]: array([[1., 0., 0.],  
          [0., 1., 0.],  
          [0., 0., 1.]])
```

```
[42]: np.diag(np.array([1,2,3,4]))
```

```
[42]: array([[1, 0, 0, 0],  
          [0, 2, 0, 0],  
          [0, 0, 3, 0],  
          [0, 0, 0, 4]])
```

NumPy dtype objects

Source: Python for Finance, 2nd ed.

Logical operations

- NumPy Arrays can be compared, just like lists.

dtype	Description	Example
?	Boolean	? (True or False)
i	Signed integer	i8 (64-bit)
u	Unsigned integer	u8 (64-bit)
f	Floating point	f8 (64-bit)
c	Complex floating point	c32 (256-bit)
m	timedelta	m (64-bit)
M	datetime	M (64-bit)
O	Object	O (pointer to object)
U	Unicode	U24 (24 Unicode characters)
V	Raw data (void)	V12 (12-byte data block)

dtype object

```
[43]: first = np.array([0, 1, 2, 3, 3, 6,])
      second = np.array([0, 1, 2, 3, 4, 5,])
```

```
[44]: first > second
```

```
[44]: array([False, False, False, False, False,  True])
```

```
[45]: first.sum() == second.sum()
```

```
[45]: True
```

```
[46]: np.any([a == 4])
```

```
[46]: True
```

```
[47]: np.all([a == 4])
```

```
[47]: False
```

Reshape and resize

- ndarray objects are immutable, but they can be reshaped (changes the view on the object) and resized (creates a new object):

```
[48]: ar = np.arange(15)
      ar
```

```
[48]: array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14])
```

```
[49]: ar.reshape((3,5))
```

```
[49]: array([[ 0,  1,  2,  3,  4],
           [ 5,  6,  7,  8,  9],
           [10, 11, 12, 13, 14]])
```

```
[50]: ar
```

```
[50]: array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14])
```

Reshape and resize

```
[51]: ar.resize((5,3))
```

```
[52]: ar
```

```
[52]: array([[ 0,  1,  2],
           [ 3,  4,  5],
           [ 6,  7,  8],
           [ 9, 10, 11],
           [12, 13, 14]])
```

Note: `reshape()` did not change the original array. `resize` did change the array's shape permanently.

Reshape and resize

- `reshape()` does not alter the total number of elements in the array.
- `resize()` can decrease (down-size) or increase (up-size) the total number of elements.

```
[53]: ar
```

```
[53]: array([[ 0,  1,  2],
           [ 3,  4,  5],
           [ 6,  7,  8],
           [ 9, 10, 11],
           [12, 13, 14]])
```

```
[54]: np.resize(ar, (3,3))
```

```
[54]: array([[0, 1, 2],
           [3, 4, 5],
           [6, 7, 8]])
```

Reshape and resize

```
[55]: np.resize(ar, (5,5))
```

```
[55]: array([[ 0,  1,  2,  3,  4],
           [ 5,  6,  7,  8,  9],
           [10, 11, 12, 13, 14],
           [ 0,  1,  2,  3,  4],
           [ 5,  6,  7,  8,  9]])
```

```
[56]: a.shape # returns the array's dimensions
```

```
[56]: (7,)
```


Further operations

- Transpose:

```
[57]: g = np.arange(0, 6)
      g.resize(2,3)
      g
```

```
[57]: array([[0, 1, 2],
            [3, 4, 5]])
```

```
[58]: g.T
```

```
[58]: array([[0, 3],
            [1, 4],
            [2, 5]])
```

- Flattening:

```
[59]: g.flatten()
```

```
[59]: array([0, 1, 2, 3, 4, 5])
```

Further operations

- Stacking: `hstack` or `vstack` can be used to connect two arrays horizontally or vertically.

```
[60]: b = np.ones((2,3))
```

```
[61]: np.vstack((g, b))
```

```
[61]: array([[0., 1., 2.],
            [3., 4., 5.],
            [1., 1., 1.],
            [1., 1., 1.]])
```

NOTE: The size of the to-be connected dimensions must be equal.

2.3 Structured NumPy arrays

Structured NumPy arrays

- The specialisation of `ndarray` may be too narrow.
- However, one can instantiate `ndarray` with a dedicated `dtype`.
- This allows to build database-like data sets where each row corresponds to an “entry”.

Structured NumPy arrays

- Creating a data type:

```
[62]: dt = np.dtype([('Name', 'S10'), ('Age', 'i4'),
                    ('Height', 'f'), ('Children/Pets', 'i4', 2)])
      dt
```

```
[62]: dtype([('Name', 'S10'), ('Age', '<i4'), ('Height', '<f4'), ('Children/Pets',
'<i4', (2,))])
```

- Equivalently:

```
[63]: dt = np.dtype({'names': ['Name', 'Age', 'Height', 'Children/Pets'],
                        'formats': '0 int float int,int'.split()})

dt
```

```
[63]: dtype([('Name', '0'), ('Age', '<i8'), ('Height', '<f8'), ('Children/Pets',
                        [('f0', '<i8'), ('f1', '<i8')])])
```

Structured NumPy arrays

- Now create the ndarray with the new data type:

```
[64]: s = np.array([('Smith', 45, 1.83, (0, 1)),
                    ('Jones', 53, 1.72, (2, 2))], dtype=dt)

s
```

```
[64]: array([('Smith', 45, 1.83, (0, 1)), ('Jones', 53, 1.72, (2, 2))],
          dtype=[('Name', '0'), ('Age', '<i8'), ('Height', '<f8'), ('Children/Pets',
                        [('f0', '<i8'), ('f1', '<i8')])])
```

```
[65]: type(s)
```

```
[65]: numpy.ndarray
```

Structured NumPy arrays

- The columns can be accessed through their names:

```
[66]: s['Name']
```

```
[66]: array(['Smith', 'Jones'], dtype=object)
```

```
[67]: s['Height'].mean()
```

```
[67]: 1.775
```

```
[68]: s[0]
```

```
[68]: ('Smith', 45, 1.83, (0, 1))
```

```
[69]: s[1]['Age']
```

```
[69]: 53
```

2.4 Data Analysis with pandas: DataFrame

Data analysis with pandas

- pandas is a powerful Python library for data manipulation and analysis. Its name is derived from panel data.
- We cover the following data structures:

Source: Python for Finance, 2nd ed.

Object type	Meaning	Used for
DataFrame	2-dimensional data object with index	Tabular data organized in columns
Series	1-dimensional data object with index	Single (time) series of data

Pandas datatypes

DataFrame Class

- `DataFrame` is a class that handles tabular data, organised in columns.
- Each row corresponds to an entry or a data record.
- It is thus similar to a table in a relational database or an Excel spreadsheet.

```
[70]: import pandas as pd

df = pd.DataFrame([10,20,30,40], # data as a list
                  columns=['numbers'], # column label
                  index=['a', 'b', 'c', 'd']) # index values for entries
```

```
[71]: df
```

```
[71]:   numbers
a      10
b      20
c      30
d      40
```

DataFrame Class

- The columns can be named (but don't need to be).
- The index can take different forms such as numbers or strings.
- The input data for the `DataFrame` Class can come in different types, such as `list`, `tuple`, `ndarray` and `dict` objects.

Simple operations

- Some simple operations applied to a `DataFrame` object:

```
[72]: df.index
```

```
[72]: Index(['a', 'b', 'c', 'd'], dtype='object')
```

```
[73]: df.columns
```

```
[73]: Index(['numbers'], dtype='object')
```

Simple operations

```
[74]: df.loc['c'] # selects value corresponding to index c
```

```
[74]: numbers    30
      Name: c, dtype: int64
```

```
[75]: df.loc[['a', 'd']] # selects values corresponding to indices a and d
```

```
[75]: numbers
a      10
d      40
```

```
[76]: df.iloc[1:3] # select second and third rows
```

```
[76]: numbers
b      20
c      30
```

Simple operations

```
[77]: df.sum()
```

```
[77]: numbers      100
dtype: int64
```

- Vectorised operations as with `ndarray`:

```
[78]: df ** 2
```

```
[78]: numbers
a      100
b      400
c      900
d     1600
```

Extending DataFrame objects

```
[79]: df['floats'] = (1.5, 2.5, 3.5, 4.5) # adds a new column
```

```
[80]: df
```

```
[80]: numbers  floats
a         10     1.5
b         20     2.5
c         30     3.5
d         40     4.5
```

```
[81]: df['floats']
```

```
[81]: a      1.5
b      2.5
c      3.5
d      4.5
Name: floats, dtype: float64
```

Extending DataFrame objects

- A `DataFrame` object can be taken to define a new column:

```
[82]: df['names'] = pd.DataFrame(['Yves', 'Sandra', 'Lilli', 'Henry'],
                                index = ['d', 'a', 'b', 'c'])
```

```
[83]: df
```

```
[83]:   numbers  floats  names
a         10     1.5  Sandra
b         20     2.5   Lilli
c         30     3.5   Henry
d         40     4.5    Yves
```

Extending DataFrame objects

- Appending data:

```
[84]: df = df.append(pd.DataFrame({'numbers': 100, 'floats': 5.75, 'names': 'Jill'},
                                index = ['y',]))
```

```
[85]: df
```

```
[85]:   numbers  floats  names
a         10     1.50  Sandra
b         20     2.50   Lilli
c         30     3.50   Henry
d         40     4.50    Yves
y        100     5.75    Jill
```

Extending DataFrame objects

- Be careful when appending without providing an index – the index gets replaced by a simple range index:

```
[86]: df.append({'numbers': 100, 'floats': 5.75, 'names': 'Jill'}, ignore_index=True)
```

```
[86]:   numbers  floats  names
0         10     1.50  Sandra
1         20     2.50   Lilli
2         30     3.50   Henry
3         40     4.50    Yves
4        100     5.75    Jill
5        100     5.75    Jill
```

Extending DataFrame objects

- Appending with missing data:

```
[87]: df = df.append(pd.DataFrame({'names': 'Liz'},
                                index = ['z']),
                                sort = False)
```

```
[88]: df
```

```
[88]:   numbers  floats  names
a      10.0     1.50  Sandra
b      20.0     2.50   Lilli
```

c	30.0	3.50	Henry
d	40.0	4.50	Yves
y	100.0	5.75	Jill
z	NaN	NaN	Liz

Mathematical operations on Data Frames

- A lot of mathematical methods are implemented for `DataFrame` objects:

```
[89]: df[['numbers', 'floats']].sum()
```

```
[89]: numbers    200.00
      floats     17.75
      dtype: float64
```

```
[90]: df['numbers'].var()
```

```
[90]: 1250.0
```

```
[91]: df['numbers'].max()
```

```
[91]: 100.0
```

Time series with Data Frame

- In this section we show how a `DataFrame` can be used to manage time series data.
- First, we create a `DataFrame` object using random numbers in an `ndarray` object.

```
[92]: import numpy as np
      import pandas as pd
      np.random.seed(100)
      a = np.random.standard_normal((9,4))
      a
```

```
[92]: array([[ -1.74976547,  0.3426804 ,  1.1530358 , -0.25243604],
             [ 0.98132079,  0.51421884,  0.22117967, -1.07004333],
             [-0.18949583,  0.25500144, -0.45802699,  0.43516349],
             [-0.58359505,  0.81684707,  0.67272081, -0.10441114],
             [-0.53128038,  1.02973269, -0.43813562, -1.11831825],
             [ 1.61898166,  1.54160517, -0.25187914, -0.84243574],
             [ 0.18451869,  0.9370822 ,  0.73100034,  1.36155613],
             [-0.32623806,  0.05567601,  0.22239961, -1.443217 ],
             [-0.75635231,  0.81645401,  0.75044476, -0.45594693]])
```

```
[93]: df = pd.DataFrame(a)
```

Note: To learn more about Python's built-in pseudo-random number generator (PRNG), see [here](#).

Practical example using DataFrame class

```
[94]: df
```

```
[94]:      0      1      2      3
0 -1.749765  0.342680  1.153036 -0.252436
1  0.981321  0.514219  0.221180 -1.070043
2 -0.189496  0.255001 -0.458027  0.435163
```

```

3 -0.583595  0.816847  0.672721 -0.104411
4 -0.531280  1.029733 -0.438136 -1.118318
5  1.618982  1.541605 -0.251879 -0.842436
6  0.184519  0.937082  0.731000  1.361556
7 -0.326238  0.055676  0.222400 -1.443217
8 -0.756352  0.816454  0.750445 -0.455947

```

Practical example using DataFrame class

- Arguments to the `DataFrame()` function for instantiating a `DataFrame` object:

Parameter	Format	Description
<code>data</code>	<code>ndarray/dict/DataFrame</code>	Data for <code>DataFrame</code> ; dict can contain <code>Series</code> , <code>ndarray</code> , <code>list</code>
<code>index</code>	<code>Index/array-like</code>	Index to use; defaults to <code>range(n)</code>
<code>columns</code>	<code>Index/array-like</code>	Column headers to use; defaults to <code>range(n)</code>
<code>dtype</code>	<code>dtype</code> , default <code>None</code>	Data type to use/force; otherwise, it is inferred
<code>copy</code>	<code>bool</code> , default <code>None</code>	Copy data from inputs

DataFrame object

Source: Python for Finance, 2nd ed.

Practical example using DataFrame class

- In the next steps, we set column names and add a time dimension for the rows.

```
[95]: df.columns = ['No1', 'No2', 'No3', 'No4']
```

```
[96]: df
```

```

[96]:
      No1      No2      No3      No4
0 -1.749765  0.342680  1.153036 -0.252436
1  0.981321  0.514219  0.221180 -1.070043
2 -0.189496  0.255001 -0.458027  0.435163
3 -0.583595  0.816847  0.672721 -0.104411
4 -0.531280  1.029733 -0.438136 -1.118318
5  1.618982  1.541605 -0.251879 -0.842436
6  0.184519  0.937082  0.731000  1.361556
7 -0.326238  0.055676  0.222400 -1.443217
8 -0.756352  0.816454  0.750445 -0.455947

```

```
[97]: df['No3'].values.flatten()
```

```

[97]: array([ 1.1530358 ,  0.22117967, -0.45802699,  0.67272081, -0.43813562,
          -0.25187914,  0.73100034,  0.22239961,  0.75044476])

```

Practical example using DataFrame class

- `pandas` is especially strong at handling times series data efficiently.
- Assume that the data rows in the `DataFrame` consist of monthly observations starting in January 2019.
- The method `date_range()` generates a `DateTimeIndex` object that can be used as the row index.

```
[98]: dates = pd.date_range('2019-1-1', periods = 9, freq = 'M')
      dates
```

```
[98]: DatetimeIndex(['2019-01-31', '2019-02-28', '2019-03-31', '2019-04-30',
                    '2019-05-31', '2019-06-30', '2019-07-31', '2019-08-31',
                    '2019-09-30'],
                    dtype='datetime64[ns]', freq='M')
```

Practical example using DataFrame class

- Parameters of the `date_range()` function:

Parameter	Format	Description
start	string/datetime	Left bound for generating dates
end	string/datetime	Right bound for generating dates
periods	integer/None	Number of periods (if start or end is None)
freq	string/DateOffset	Frequency string, e.g., 5D for 5 days
tz	string/None	Time zone name for localized index
normalize	bool, default None	Normalizes start and end to midnight
name	string, default None	Name of resulting index

Date range parameters

Source: Python for Finance, 2nd ed.

Practical example using DataFrame class

- Frequency parameter of `date_range()` function:

```
<table>
  <td></td>
  <td></td>
</table>
```

Source: Python for Finance, 2nd ed.

Practical example using DataFrame class

- Now set the row index to the dates:

```
[99]: df.index = dates

df
```

```
[99]:
```

	No1	No2	No3	No4
2019-01-31	-1.749765	0.342680	1.153036	-0.252436
2019-02-28	0.981321	0.514219	0.221180	-1.070043
2019-03-31	-0.189496	0.255001	-0.458027	0.435163
2019-04-30	-0.583595	0.816847	0.672721	-0.104411
2019-05-31	-0.531280	1.029733	-0.438136	-1.118318
2019-06-30	1.618982	1.541605	-0.251879	-0.842436
2019-07-31	0.184519	0.937082	0.731000	1.361556


```
2019-08-31 -0.326238  0.055676  0.222400 -1.443217
2019-09-30 -0.756352  0.816454  0.750445 -0.455947
```

Practical example using DataFrame class

- Next, we visualise the data:

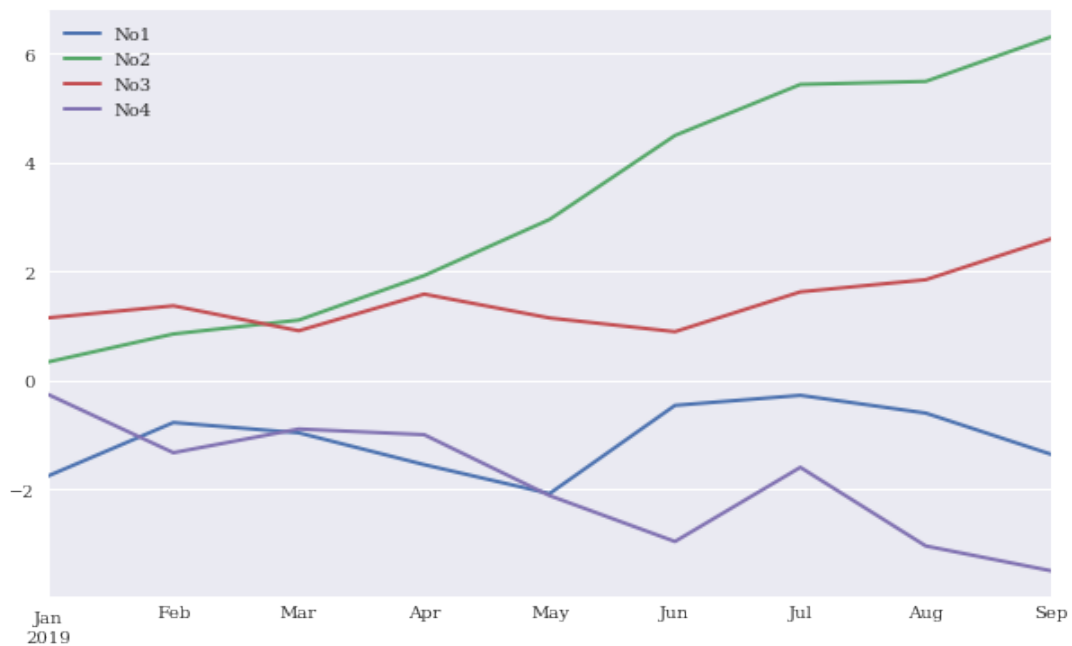
```
[100]: from pylab import plt, mpl # imports for visualisation
plt.style.use('seaborn') # This and the following lines customise the plot style
mpl.rcParams['font.family'] = 'serif'
%matplotlib inline
```

- More about customising the plot style: [here](#).

Practical example using DataFrame class

- Plot the cumulative sum for each column of df:

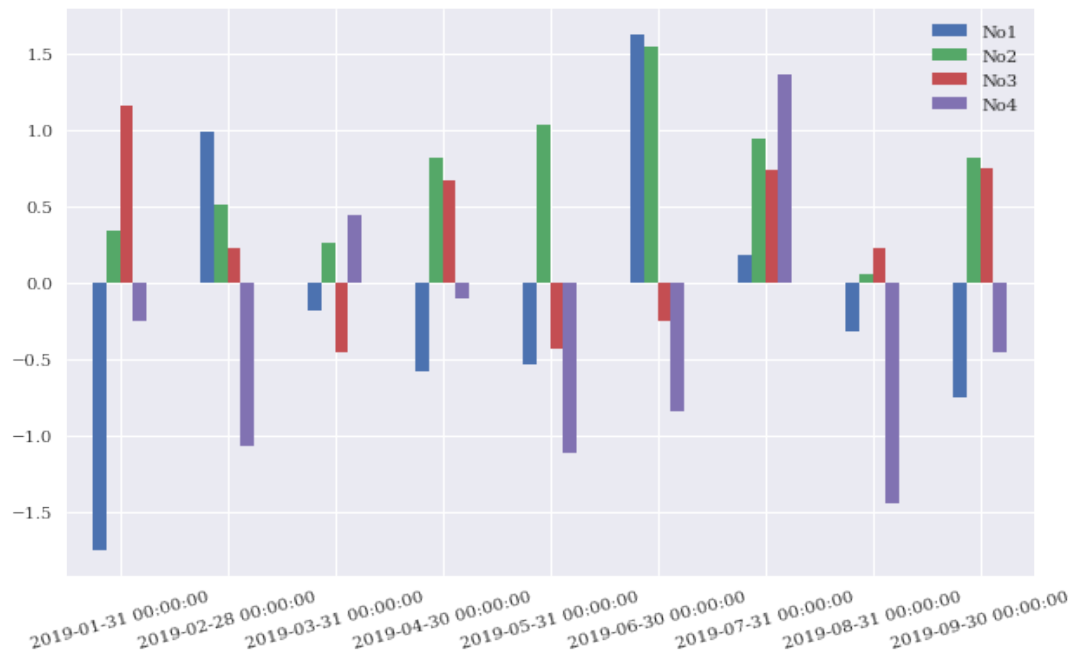
```
[101]: df.cumsum().plot(lw = 2.0, figsize = (10,6));
```



Practical example using DataFrame class

- A bar chart:

```
[102]: df.plot.bar(figsize = (10,6), rot = 15);
```



Practical example using DataFrame class

- Parameters of plot() method:

Parameter	Format	Description
x	label/position, default None	Only used when column values are x-ticks
y	label/position, default None	Only used when column values are y-ticks
subplots	boolean, default False	Plot columns in subplots
sharex	boolean, default True	Share the x-axis
sharey	boolean, default False	Share the y-axis
use_index	boolean, default True	Use DataFrame.index as x-ticks
stacked	boolean, default False	Stack (only for bar plots)
sort_columns	boolean, default False	Sort columns alphabetically before plotting
title	string, default None	Title for the plot
grid	boolean, default False	Show horizontal and vertical grid lines
legend	boolean, default True	Show legend of labels
ax	matplotlib axis object	matplotlib axis object to use for plotting
style	string or list/dictionary	Line plotting style (for each column)
kind	string (e.g., "line", "bar", "barh", "kde", "density")	Type of plot
logx	boolean, default False	Use logarithmic scaling of x-axis
logy	boolean, default False	Use logarithmic scaling of y-axis
xticks	sequence, default Index	X-ticks for the plot

Parameters of plot method

Source: Python for Finance, 2nd ed.

Practical example using DataFrame class

- Parameters of plot() method:

```
<table>
  <td align=top></td>
</table>
```

Source: Python for Finance, 2nd ed.

Practical example using DataFrame class

- Useful functions:

```
[103]: df.info() # provide basic information
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 9 entries, 2019-01-31 to 2019-09-30
Freq: M
Data columns (total 4 columns):
#   Column  Non-Null Count  Dtype
---  -
0    No1      9 non-null        float64
1    No2      9 non-null        float64
2    No3      9 non-null        float64
3    No4      9 non-null        float64
dtypes: float64(4)
memory usage: 360.0 bytes
```

Practical example using DataFrame class

```
[104]: df.sum()
```

```
[104]: No1    -1.351906
      No2     6.309298
      No3     2.602739
      No4    -3.490089
      dtype: float64
```

```
[105]: df.mean(axis=0) # column-wise mean
```

```
[105]: No1    -0.150212
      No2     0.701033
      No3     0.289193
      No4    -0.387788
      dtype: float64
```

```
[106]: df.mean(axis=1) # row-wise mean
```

```
[106]: 2019-01-31    -0.126621
      2019-02-28     0.161669
      2019-03-31     0.010661
      2019-04-30     0.200390
      2019-05-31    -0.264500
      2019-06-30     0.516568
      2019-07-31     0.803539
      2019-08-31    -0.372845
      2019-09-30     0.088650
      Freq: M, dtype: float64
```

Useful functions: `groupby()`

```
[107]: df['Quarter'] = ['Q1', 'Q1', 'Q1', 'Q2', 'Q2', 'Q2', 'Q3', 'Q3', 'Q3',]
```

```
[108]: df
```

```
[108]:
```

	No1	No2	No3	No4	Quarter
2019-01-31	-1.749765	0.342680	1.153036	-0.252436	Q1
2019-02-28	0.981321	0.514219	0.221180	-1.070043	Q1
2019-03-31	-0.189496	0.255001	-0.458027	0.435163	Q1
2019-04-30	-0.583595	0.816847	0.672721	-0.104411	Q2
2019-05-31	-0.531280	1.029733	-0.438136	-1.118318	Q2
2019-06-30	1.618982	1.541605	-0.251879	-0.842436	Q2
2019-07-31	0.184519	0.937082	0.731000	1.361556	Q3
2019-08-31	-0.326238	0.055676	0.222400	-1.443217	Q3
2019-09-30	-0.756352	0.816454	0.750445	-0.455947	Q3

Useful functions: `groupby()`

```
[109]: groups = df.groupby('Quarter')
```

```
[110]: groups.mean()
```

```
[110]:
```

	No1	No2	No3	No4
Quarter				
Q1	-0.319314	0.370634	0.305396	-0.295772
Q2	0.168035	1.129395	-0.005765	-0.688388
Q3	-0.299357	0.603071	0.567948	-0.179203

```
[111]: groups.max()
```

```
[111]:
```

	No1	No2	No3	No4
Quarter				
Q1	0.981321	0.514219	1.153036	0.435163
Q2	1.618982	1.541605	0.672721	-0.104411
Q3	0.184519	0.937082	0.750445	1.361556

Useful functions: `groupby()`

```
[112]: groups.aggregate([min, max]).round(3)
```

```
[112]:
```

	No1		No2		No3		No4	
	min	max	min	max	min	max	min	max
Quarter								
Q1	-1.750	0.981	0.255	0.514	-0.458	1.153	-1.070	0.435
Q2	-0.584	1.619	0.817	1.542	-0.438	0.673	-1.118	-0.104
Q3	-0.756	0.185	0.056	0.937	0.222	0.750	-1.443	1.362

Selecting and filtering data

- Logical operators can be used to filter data.
- First, construct a `DataFrame` filled with random numbers to work with.

```
[113]: data = np.random.standard_normal((10,2))
```

```
[114]: df = pd.DataFrame(data, columns = ['x', 'y'])
```

```
[115]: df.head(2) # the first two rows
```

```
[115]:
```

	x	y
0	1.189622	-1.690617
1	-1.356399	-1.232435

```
[116]: df.tail(2) # the last two rows
```

```
[116]:
```

	x	y
8	-0.940046	-0.827932
9	0.108863	0.507810

Selecting and filtering data

```
[117]: (df['x'] > 1) & (df['y'] < 1) # check if value in x-column is greater than 1 and  
      ↪ value in y-column is smaller than 1
```

```
[117]:
```

0	True
1	False
2	False
3	False
4	True
5	False
6	False
7	False
8	False
9	False

dtype: bool

```
[118]: df[df['x'] > 1]
```

```
[118]:
```

	x	y
0	1.189622	-1.690617
4	1.299748	-1.733096

```
[119]: df.query('x > 1') # query()-method takes string as parameter
```

```
[119]:
```

	x	y
0	1.189622	-1.690617
4	1.299748	-1.733096

Selecting and filtering data

```
[120]: (df > 1).head(3) # Find values greater than 1
```

```
[120]:
```

	x	y
0	True	False
1	False	False
2	False	False

```
[121]: df[df > 1].head(3) # Select values greater than 1 and put NaN (not-a-number) in the  
      ↪ other entries
```

```
[121]:
```

	x	y
0	1.189622	NaN
1	NaN	NaN
2	NaN	NaN

Concatenation

- Adding rows from one data frame to another data frame can be done with `append()` or `concat()`:

```
[122]: df1 = pd.DataFrame(['100', '200', '300', '400'],
                        index = ['a', 'b', 'c', 'd'],
                        columns = ['A',])

df2 = pd.DataFrame(['200', '150', '50'],
                    index = ['f', 'b', 'd'],
                    columns = ['B',])
```

Concatenation

```
[123]: df1.append(df2, sort = False)
```

```
[123]:
```

	A	B
a	100	NaN
b	200	NaN
c	300	NaN
d	400	NaN
f	NaN	200
b	NaN	150
d	NaN	50

Concatenation

```
[124]: pd.concat((df1, df2), sort = False)
```

```
[124]:
```

	A	B
a	100	NaN
b	200	NaN
c	300	NaN
d	400	NaN
f	NaN	200
b	NaN	150
d	NaN	50

Joining

- In Python, `join()` refers to joining `DataFrame` objects according to their index values.
- There are four different types of joining:
 1. left join
 2. right join
 3. inner join
 4. outer join

Joining

```
[125]: df1.join(df2, how = 'left') # default join, based on indices of first dataset
```

```
[125]:
```

	A	B
a	100	NaN
b	200	150
c	300	NaN
d	400	50

```
[126]: df1.join(df2, how = 'right') # based on indices of second dataset
```

```
[126]:
```

	A	B
f	NaN	200
b	200	150
d	400	50

Joining

```
[127]: df1.join(df2, how = 'inner') # preserves those index values that are found in both ↵  
      ↪ datasets
```

```
[127]:
```

	A	B
b	200	150
d	400	50

```
[128]: df1.join(df2, how = 'outer') # preserves indices found in both datasets
```

```
[128]:
```

	A	B
a	100	NaN
b	200	150
c	300	NaN
d	400	50
f	NaN	200

Merging

- Join operations on `DataFrame` objects are based on the datasets indices.
- **Merging** operates on a shared column of two `DataFrame` objects.
- To demonstrate the usage we add a new column `C` to `df1` and `df2`.

```
[129]: c = pd.Series([250, 150, 50], index = ['b', 'd', 'c'])  
      df1['C'] = c  
      df2['C'] = c
```

Merging

```
[130]: df1
```

```
[130]:
```

	A	C
a	100	NaN
b	200	250.0
c	300	50.0
d	400	150.0

```
[131]: df2
```

```
[131]:
```

	B	C
f	200	NaN
b	150	250.0
d	50	150.0

Merging

- By default, a merge takes place on a shared column, preserving only the shared data rows:

```
[132]: pd.merge(df1, df2)
```

```
[132]:
```

	A	C	B
0	100	NaN	200
1	200	250.0	150
2	400	150.0	50

- An **outer merge** preserves all data rows:

```
[133]: pd.merge(df1, df2, how = 'outer')
```

```
[133]:
```

	A	C	B
0	100	NaN	200
1	200	250.0	150
2	300	50.0	NaN
3	400	150.0	50

Merging

- There are numerous other ways to merge `DataFrame` objects.
- To learn more about merging in Python, see the pandas document on [DataFrame merging](#).

```
[134]: pd.merge(df1, df2, left_on = 'A', right_on = 'B')
```

```
[134]:
```

	A	C_x	B	C_y
0	200	250.0	200	NaN

```
[135]: pd.merge(df1, df2, left_on = 'A', right_on = 'B', how = 'outer')
```

```
[135]:
```

	A	C_x	B	C_y
0	100	NaN	NaN	NaN
1	200	250.0	200	NaN
2	300	50.0	NaN	NaN
3	400	150.0	NaN	NaN
4	NaN	NaN	150	250.0
5	NaN	NaN	50	150.0